

# Elf Loader

March 14, 2025

## 1 Introduction

The program in this section serves two purposes:

- Understanding ELF binaries: By compiling and analyzing a simple C program (`hello.c`), we will extract and examine the ELF header, section header table, program header table, and various other components.
- Demonstrating `binfmt`: By compiling Python scripts into `.pyc` files, we explore how they can be executed using `binfmt_misc`, allowing non-native binaries to run seamlessly.

## 2 Background

The Executable and Linkable Format (ELF) is the standard binary format used in Unix-like operating systems, including Linux, for executables, object files, shared libraries, and core dumps. An ELF loader is a critical component of the operating system responsible for loading these binaries into memory and preparing them for execution.

## 3 Programs

### 3.1 PART 1: Analyzing ELF Layout of Binaries

#### 3.1.1 Compiling the C Program

To generate an ELF binary, we compile `hello.c` using:

```
gcc -o hello hello.c
```

#### 3.1.2 Viewing the ELF Header

The ELF header provides essential metadata about the binary.

```
readelf -h hello
```

```

ELF Header:
Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
Class:   ELF64
Data:    2's complement, little endian
Version: 1 (current)
OS/ABI:  UNIX - System V
ABI Version: 0
Type:    EXEC (Executable file)
Machine: Advanced Micro Devices X86-64
Version: 0x1
Entry point address: 0x400410
Start of program headers: 64 (bytes into file)
Start of section headers: 15976 (bytes into file)
Flags:   0x0
Size of this header: 64 (bytes)
Size of program headers: 56 (bytes)
Number of program headers: 9
Size of section headers: 64 (bytes)
Number of section headers: 30
Section header string table index: 29

```

Figure 1: ELF Header: Shows architecture type, entry point, and ELF class (32-bit or 64-bit). (Fig 1).

The ELF Header provides crucial metadata about the binary format, helping the OS loader understand how to process the executable. Below is a breakdown of the key fields:

- **Magic Number (7f 45 4c 46):** Identifies the file as an ELF binary. The first byte (0x7f) is a marker, followed by "ELF" in ASCII (45 4c 46).
- **Class (ELF64):** Indicates that this is a 64-bit ELF binary.
- **Data (2's complement, little endian):** Specifies the byte order (Little Endian format).
- **Version (1 - current):** ELF format version.
- **OS/ABI (UNIX - System V):** Target operating system. System V is the standard ABI used on Linux.
- **Type (EXEC - Executable file):** Denotes this as an executable binary.
- **Machine (Advanced Micro Devices X86-64):** Target architecture (AMD x86-64).
- **Entry Point Address (0x400410):** The memory address where execution begins when the program is loaded.
- **Start of Program Headers (64 bytes into file):** The offset where the program header table begins.
- **Start of Section Headers (15976 bytes into file):** The offset where the section header table begins.
- **Size of ELF Header (64 bytes):** The fixed size of the ELF header.

- **Size of Program Headers (56 bytes)**
- **Number of Program Headers (9)**
- **Size of Section Headers (64 bytes)**
- **Number of Section Headers (30)**
- **Section Header String Table Index (29)**

This header information is critical for the loader.

### 3.1.3 Viewing the Section Header Table

Lists all sections within the ELF binary.

```
readelf -S hello
```

The **Section Header Table** provides a structured view of the different sections that make up the ELF binary. Each section serves a specific role in the program's execution and linking process. Below is a breakdown of the key sections:

- **Metadata and Linking Sections**
  - `.interp` - Contains the program interpreter path.
  - `.note.ABI-tag` - Stores ABI-related metadata.
  - `.gnu.hash`, `.hash` - Used for symbol lookup in dynamic linking.
  - `.dynsym`, `.dynstr`, `.gnu.version` - Store symbol tables for dynamic linking.
- **Executable Code Sections**
  - `.text` - Contains the main program's executable code.
  - `.init`, `.fini` - Used for initialization and termination routines.
  - `.plt` - Holds the procedure linkage table for dynamic function calls.
- **Data and Read-Only Sections**
  - `.rodata` - Stores read-only data such as string literals.
  - `.data` - Stores initialized global and static variables.
  - `.bss` - Reserves space for uninitialized global variables.
- **Symbol and String Tables**
  - `.symtab` - Contains symbol information (not always present in stripped binaries).
  - `.strtab` - Holds string data for symbol names.
  - `.shstrtab` - Stores section names.

### 3.1.4 Viewing the Program Header Table

Displays segments loaded into memory.

```
readelf -l hello
```

The Program Header Table describes how the OS loader should map the executable into memory. It defines segments, their locations, sizes, and permissions. This ELF file has 9 program headers, each serving a specific role in execution.

- **PHDR (Program Header Table)** - Defines the location of the program header table itself, helping the ELF loader find other headers.
- **INTERP (Interpreter Path)** - Specifies the dynamic linker used to load shared libraries. In this case, it is `/lib64/ld-linux-x86-64.so.2`.
- **LOAD (Code and Data Segments)**
  - First LOAD segment (RX) - Contains executable code (`.text` section).
  - Second LOAD segment (RW) - Stores writable data (`.data`, `.bss`).
- **DYNAMIC (Dynamic Linking Information)** - Holds dynamic linking data required at runtime.
- **NOTE (Metadata Section)** - Stores metadata such as ABI information.
- **GNU\_EH\_FRAME (Exception Handling)** - Contains data for stack unwinding during exceptions.
- **GNU\_STACK (Stack Permissions)** - Defines stack properties, marked as **Read/Write (RW)** but **not executable (NX)** for security.
- **GNU\_RELRO (Relocation Read-Only)** - Marks certain sections as **read-only after relocation** to protect against exploits.

### 3.1.5 Sections and Segments Mapping

Shows how sections map to segments.

```
readelf -l hello
```

The section-to-segment mapping shows how sections (logical units in an ELF file) are grouped into segments (units mapped into memory by the OS loader).

### 3.1.6 Extracting the Symbol Table

Displays function and variable symbols.

```
readelf -s hello
```

The **symbol table** in an ELF binary provides a list of all symbols (functions, variables, and special markers) used in the program. The key symbol tables in this output are:

- **.dynsym (Dynamic Symbol Table)**
  - Contains **symbols required for dynamic linking**.
  - Includes external functions like `printf` and `__libc_start_main`, which are linked dynamically from **GLIBC**.
  - `__gmon_start__` is a weak symbol often related to **profiling/debugging**.
- **.symtab (Full Symbol Table)**
  - Contains **all symbols** in the executable, including:
    - \* **Functions:** `main`, `_start`, `deregister_tm_clones`, etc.
    - \* **Global Variables:** `_edata`, `_bss_start`, `_end`.
    - \* **Sections and markers:** `.init`, `.plt`, `.text`.
  - Includes **file references** such as `hello.c` and `crtstuff.c`, indicating where symbols originated.

### 3.1.7 Extracting the String Table

Lists string literals stored in the binary.

```
readelf -p .strtab hello
```

The String Table in an ELF binary stores symbol names, section names, and dynamic linking information in a compact format. Instead of storing full names in the symbol table, ELF uses offsets into the string table to save space. The section string table names sections, while the symbol string tables name functions and variables.

### 3.1.8 Viewing Relocation Entries

Lists relocations used for dynamic linking.

```
readelf -r hello
```

The relocation table in an ELF binary contains entries that help resolve symbol addresses at runtime. These entries modify specific memory locations to contain correct addresses before the program executes.

- **.rela.dyn (Relocation for Dynamic Linking)**
  - Used for **global symbols** that must be resolved at runtime.
  - Entries:

- \* `__libc_start_main@GLIBC_2.2.5` - The entry point of the C runtime.
- \* `__gmon_start_` - Used for profiling/debugging.
- **.rela.plt (Relocation for Procedure Linkage Table - PLT)**
  - Used for dynamically linked functions in shared libraries.
  - Entry:
    - \* `printf@GLIBC_2.2.5` - This function call needs to be resolved at runtime.

## 3.2 PART 2: Demonstrating binfmt with Compiled Python Program

### 3.2.1 Compiling the Python Script

We compile the Python script into bytecode using:

```
python -m py_compile example.py
```

### 3.2.2 Registering Python Bytecode with binfmt\_misc

To enable direct execution of compiled Python files, we register them in binfmt:

```
sudo mount binfmt_misc -t binfmt_misc /proc/sys/fs/binfmt_misc
echo ':python_magic:E::\x03\xfb\x0d\x0a::/usr/bin/python3:' | sudo tee /proc/sys
```

### 3.2.3 Making Python Bytecode Executable

After registration, we make `example.pyc` executable:

```
chmod +x example.pyc
./example.pyc
```

## 4 Takeaway

ELF Analysis:

- Read ELF metadata using `readelf` and `objdump`.
- Understand headers, segments, symbols, and relocation.

`binfmt_misc`:

- Enables execution of compiled Python bytecode without explicitly invoking `python`.
- Registers a magic number to map `.pyc` files to Python.

There are 30 section headers, starting at offset 0x3e68:

Section Headers:									
[Nr]	Name	Type	Address	Offset	Flags	Link	Info	Align	
	Size	EntSize							
[ 0]	0000000000000000	NULL	0000000000000000	00000000	0	0	0	0	
[ 1]	.interp 000000000000001c	PROGBITS	0000000000400238	00000238	A	0	0	1	
[ 2]	.note.ABI-tag 0000000000000020	NOTE	0000000000400254	00000254	A	0	0	4	
[ 3]	.hash 0000000000000024	HASH	0000000000400278	00000278	A	5	0	8	
[ 4]	.gnu.hash 000000000000001c	GNU_HASH	00000000004002a0	000002a0	A	5	0	8	
[ 5]	.dynsym 0000000000000060	DYNSYM	00000000004002c0	000002c0	A	6	1	8	
[ 6]	.dynstr 000000000000003f	STRTAB	0000000000400320	00000320	A	0	0	1	
[ 7]	.gnu.version 0000000000000008	VERSYM	0000000000400360	00000360	A	5	0	2	
[ 8]	.gnu.version_r 0000000000000020	VERNEED	0000000000400368	00000368	A	6	1	8	
[ 9]	.rela.dyn 0000000000000030	RELA	0000000000400388	00000388	A	5	0	8	
[10]	.rela.plt 0000000000000018	RELA	00000000004003b8	000003b8	AI	5	22	8	
[11]	.init 000000000000001b	PROGBITS	00000000004003d0	000003d0	AX	0	0	4	
[12]	.plt 0000000000000020	PROGBITS	00000000004003f0	000003f0	AX	0	0	16	
[13]	.text 00000000000000175	PROGBITS	0000000000400410	00000410	AX	0	0	16	
[14]	.fini 000000000000000d	PROGBITS	0000000000400588	00000588	AX	0	0	4	
[15]	.rodata 0000000000000010	PROGBITS	0000000000400598	00000598	A	0	0	4	
[16]	.eh_frame_hdr 000000000000003c	PROGBITS	00000000004005a8	000005a8	A	0	0	4	
[17]	.eh_frame 00000000000000e8	PROGBITS	00000000004005e8	000005e8	A	0	0	8	
[18]	.init_array 0000000000000008	INIT_ARRAY	0000000000600e00	00000e00	WA	0	0	8	
[19]	.fini_array 0000000000000008	FINI_ARRAY	0000000000600e08	00000e08	WA	0	0	8	
[20]	.dynamic 000000000000001e0	DYNAMIC	0000000000600e10	00000e10	WA	6	0	8	
[21]	.got 0000000000000010	PROGBITS	0000000000600ff0	00000ff0	WA	0	0	8	
[22]	.got.plt 0000000000000020	PROGBITS	0000000000601000	00001000	WA	0	0	8	
[23]	.data 0000000000000010	PROGBITS	0000000000601020	00001020	WA	0	0	8	
[24]	.bss 0000000000000008	NOBITS	0000000000601030	00001030	WA	0	0	1	
[25]	.comment 000000000000003e	PROGBITS	0000000000000000	00001030	MS	0	0	1	
[26]	.gnu.build.attrib 000000000000001dac	NOTE	0000000000a01038	00001070		0	0	4	
[27]	.symtab	SYMTAB	0000000000000000	00002e20					

Figure 2: Section Header Table: Shows sections such as .text, .data, .bss. Key to Flags: W (write), A (alloc), X (execute), M (merge), S (strings), I (info), L (link order), O (extra OS processing required), G (group), T (TLS), C (compressed), x (unknown), o (OS specific), E (exclude), l (large), p (processor specific)

```

Elf file type is EXEC (Executable file)
Entry point 0x400410
There are 9 program headers, starting at offset 64

Program Headers:
Type           Offset             VirtAddr           PhysAddr
               FileSiz            MemSiz             Flags  Align
PHDR           0x0000000000000040 0x0000000000400040 0x0000000000400040
               0x0000000000001f8 0x0000000000001f8  R      0x8
INTERP         0x0000000000000238 0x0000000000400238 0x0000000000400238
               0x00000000000001c 0x00000000000001c  R      0x1
                [Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]
LOAD           0x0000000000000000 0x0000000000400000 0x0000000000400000
               0x00000000000006d0 0x00000000000006d0  R E    0x200000
LOAD           0x0000000000000e00 0x0000000000000e00 0x0000000000000e00
               0x0000000000000230 0x0000000000000230  RW     0x200000
DYNAMIC        0x0000000000000e10 0x0000000000000e10 0x0000000000000e10
               0x00000000000001e0 0x00000000000001e0  RW     0x8
NOTE           0x0000000000000254 0x0000000000400254 0x0000000000400254
               0x0000000000000020 0x0000000000000020  R      0x4
GNU_EH_FRAME   0x00000000000005a8 0x00000000004005a8 0x00000000004005a8
               0x000000000000003c 0x000000000000003c  R      0x4
GNU_STACK      0x0000000000000000 0x0000000000000000 0x0000000000000000
               0x0000000000000000 0x0000000000000000  RW     0x10
GNU_RELRO      0x0000000000000e00 0x0000000000000e00 0x0000000000000e00
               0x0000000000000200 0x0000000000000200  R      0x1

Section to Segment mapping:
Segment Sections...
00
01      .interp
02      .interp .note.ABI-tag .hash .gnu.hash .dynsym .dynstr .gnu.version .gnu.ver
r .eh_frame
03      .init_array .fini_array .dynamic .got .got.plt .data .bss
04      .dynamic
05      .note.ABI-tag
06      .eh_frame_hdr
07
08      .init_array .fini_array .dynamic .got

```

Figure 3: Program Header Table: Lists executable segments and memory protection attributes.

```

Section to Segment mapping:
Segment Sections...
00
01      .interp
02      .interp .note.ABI-tag .hash .gnu.hash .dynsym .dynstr .gnu.version .gnu.version_r .rela.dyn .rela.plt .init .p
lt .text .fini .rodata .eh_frame_hdr .eh_frame
03      .init_array .fini_array .dynamic .got .got.plt .data .bss
04      .dynamic
05      .note.ABI-tag
06      .eh_frame_hdr
07
08      .init_array .fini_array .dynamic .got

```

Figure 4: Mapping between sections and segments in ELF.



```

Symbol table '.dynsym' contains 4 entries:
Num:  Value      Size Type  Bind  Vis  Ndx Name
  0: 0000000000000000 0 NOTYPE LOCAL DEFAULT UND
  1: 0000000000000000 0 FUNC  GLOBAL DEFAULT UND printf@GLIBC_2.2.5 (2)
  2: 0000000000000000 0 FUNC  GLOBAL DEFAULT UND __libc_start_main@GLIBC_2.2.5 (2)
  3: 0000000000000000 0 NOTYPE WEAK  DEFAULT UND __gmon_start__

Symbol table '.symtab' contains 100 entries:
Num:  Value      Size Type  Bind  Vis  Ndx Name
  0: 0000000000000000 0 NOTYPE LOCAL DEFAULT UND
  1: 0000000000400238 0 SECTION LOCAL DEFAULT 1
  2: 0000000000400254 0 SECTION LOCAL DEFAULT 2
  3: 0000000000400278 0 SECTION LOCAL DEFAULT 3
  4: 00000000004002a0 0 SECTION LOCAL DEFAULT 4
  5: 00000000004002c0 0 SECTION LOCAL DEFAULT 5
  6: 0000000000400320 0 SECTION LOCAL DEFAULT 6
  7: 0000000000400360 0 SECTION LOCAL DEFAULT 7
  8: 0000000000400368 0 SECTION LOCAL DEFAULT 8
  9: 0000000000400388 0 SECTION LOCAL DEFAULT 9
 10: 00000000004003b8 0 SECTION LOCAL DEFAULT 10
 11: 00000000004003d0 0 SECTION LOCAL DEFAULT 11
 12: 00000000004003f0 0 SECTION LOCAL DEFAULT 12
 13: 0000000000400410 0 SECTION LOCAL DEFAULT 13
 14: 0000000000400588 0 SECTION LOCAL DEFAULT 14
 15: 0000000000400598 0 SECTION LOCAL DEFAULT 15
 16: 00000000004005a8 0 SECTION LOCAL DEFAULT 16
 17: 00000000004005e8 0 SECTION LOCAL DEFAULT 17
 18: 0000000000600e00 0 SECTION LOCAL DEFAULT 18
 19: 0000000000600e08 0 SECTION LOCAL DEFAULT 19
 20: 0000000000600e10 0 SECTION LOCAL DEFAULT 20
 21: 0000000000600ff0 0 SECTION LOCAL DEFAULT 21
 22: 0000000000601000 0 SECTION LOCAL DEFAULT 22
 23: 0000000000601020 0 SECTION LOCAL DEFAULT 23
 24: 0000000000601030 0 SECTION LOCAL DEFAULT 24
 25: 0000000000000000 0 SECTION LOCAL DEFAULT 25
 26: 0000000000a01038 0 SECTION LOCAL DEFAULT 26
 27: 0000000000000000 0 FILE  LOCAL DEFAULT ABS /lib/../lib64/crt1.o
 28: 000000000040043f 0 NOTYPE LOCAL HIDDEN 13 .annobin_init.c
 29: 000000000040043f 0 NOTYPE LOCAL HIDDEN 13 .annobin_init.c_end
 30: 0000000000400410 0 NOTYPE LOCAL HIDDEN 13 .annobin_init.c.hot

```

Figure 5: Symbol Table: Shows function symbols, global variables, and linkage information.

```

[ 416] crtstuff.c
[ 421] deregister_tm_clones
[ 436] __do_global_dtors_aux
[ 44c] completed.7312
[ 45b] __do_global_dtors_aux_fini_array_entry
[ 482] frame_dummy
[ 48e] __frame_dummy_init_array_entry
[ 4ad] hello.c
[ 4b5] __FRAME_END__
[ 4c3] __init_array_end
[ 4d4] _DYNAMIC
[ 4dd] __init_array_start
[ 4f0] __GNU_EH_FRAME_HDR
[ 503] _GLOBAL_OFFSET_TABLE_
[ 519] __libc_csu_fini
[ 529] _edata
[ 530] printf@@GLIBC_2.2.5
[ 544] __libc_start_main@@GLIBC_2.2.5
[ 563] __data_start
[ 570] __gmon_start__
[ 57f] __dso_handle
[ 58c] _IO_stdin_used
[ 59b] __libc_csu_init
[ 5ab] _dl_relocate_static_pie
[ 5c3] __bss_start
[ 5cf] main
[ 5d4] __TMC_END__

```

Figure 6: String Table: Displays strings embedded within the ELF binary.

```

Relocation section '.rela.dyn' at offset 0x388 contains 2 entries:
  Offset      Info            Type             Sym. Value      Sym. Name + Addend
0000000600ff0 00020000000006 R_X86_64_GLOB_DAT 0000000000000000 __libc_start_main@GLIBC_2.2.5 + 0
0000000600ff8 00030000000006 R_X86_64_GLOB_DAT 0000000000000000 __gmon_start__ + 0

Relocation section '.rela.plt' at offset 0x3b8 contains 1 entry:
  Offset      Info            Type             Sym. Value      Sym. Name + Addend
0000000601018 00010000000007 R_X86_64_JUMP_SLO 0000000000000000 printf@GLIBC_2.2.5 + 0

```

Figure 7: Relocation Table: Shows addresses modified at runtime.

```

[~/operating-systems-course/elf_loaders]$ sudo mount binfmt_misc -t binfmt_misc /proc/sys/fs/binfmt_misc
[sudo] password for [redacted]:
mount: /proc/sys/fs/binfmt_misc: binfmt_misc already mounted on /proc/sys/fs/binfmt_misc.
[~/operating-systems-course/elf_loaders]$ cat /proc/sys/fs/binfmt_misc/python3.10
enabled
interpreter /usr/bin/python3.10
flags:
offset 0
magic 0f0d0d0a

```

Figure 8: Mount.